

Overview

This notebook will show you how to create and query a table or DataFrame that you uploaded to DBFS. [DBFS \(https://docs.databricks.com/user-guide/dbfs-databricks-file-system.html\)](https://docs.databricks.com/user-guide/dbfs-databricks-file-system.html) is a Databricks File System that allows you to store data for querying inside of Databricks. This notebook assumes that you have a file already inside of DBFS that you would like to read from.

This notebook is written in **Python** so the default cell type is Python. However, you can use different languages by using the `%LANGUAGE` syntax. Python, Scala, SQL, and R are all supported.

In [2]:

```
# Importação Bibliotecas
from pyspark.sql import SparkSession
from pyspark import HiveContext
from pyspark.sql.functions import monotonically_increasing_id
from pyspark.ml.linalg import Vectors
from pyspark.ml.feature import VectorAssembler
from pyspark.ml.evaluation import RegressionEvaluator
from pyspark.ml.classification import LogisticRegression, LogisticRegressionModel
from pyspark.mllib.evaluation import BinaryClassificationMetrics as metric
from pyspark.ml.feature import OneHotEncoder, StringIndexer, VectorAssembler, OneHotEncoder
from pyspark.ml.classification import RandomForestClassifier, RandomForestClassificationMetrics
from pyspark.mllib.evaluation import BinaryClassificationMetrics as metric
from pyspark.ml import Pipeline
```

In [3]:

```
spark.sparkContext._conf.getAll()
```

Out[3]:

```
[('spark.sql.catalogImplementation', 'hive'),
 ('spark.rdd.compress', 'True'),
 ('spark.driver.host', '10.30.30.21'),
 ('spark.app.id', 'local-1559699171437'),
 ('spark.serializer.objectStreamReset', '100'),
 ('spark.master', 'local[*]'),
 ('spark.executor.id', 'driver'),
 ('spark.submit.deployMode', 'client'),
 ('spark.app.name', 'PySparkShell'),
 ('spark.driver.port', '35371')]
```

In [4]:

```
conf = spark.sparkContext._conf.setAll([
    ("hive.metastore.uris", "thrift://localhost:9083")])
```

In [5]:

```
spark.stop()
```

In [6]:

```
sc = SparkContext()
```

In [7]:

```
spark.sparkContext._conf.getAll()
```

Out[7]:

```
[('spark.sql.catalogImplementation', 'hive'),
 ('spark.rdd.compress', 'True'),
 ('spark.driver.host', '10.30.30.21'),
 ('spark.app.id', 'local-1559699171437'),
 ('hive.metastore.uris', 'thrift://localhost:9083'),
 ('spark.serializer.objectStreamReset', '100'),
 ('spark.master', 'local[*]'),
 ('spark.executor.id', 'driver'),
 ('spark.submit.deployMode', 'client'),
 ('spark.app.name', 'PySparkShell'),
 ('spark.driver.port', '35371')]
```

In [8]:

```
df = spark.sql("SHOW TABLES")
df.show()
```

```
+-----+-----+-----+
|database|      tableName|isTemporary|
+-----+-----+-----+
| default|boosting_output|      false|
+-----+-----+-----+
```

In [9]:

```
spark = SparkSession.builder.config(conf=conf).getOrCreate()
```

In [10]:

```
# Carga do Arquivo via HDFS
#file_location = "/FileStore/tables/bank_additional_full_no_header-a99ad.csv"
file_location = 'hdfs:///user/labdata/marketing/bank-additional-full.csv'
file_type = "csv"

# CSV options
infer_schema = "true"
first_row_is_header = "false"
delimiter = ";"

# The applied options are for CSV files. For other file types, these will be ignored.
df_marketing_data = spark.read.format(file_type) \
    .option("inferSchema", infer_schema) \
    .option("header", first_row_is_header) \
    .option("sep", delimiter) \
    .load(file_location)
```

In [11]:

```
# Definição dos Nomes de Variaveis
DefColumnNames=df_marketing_data.schema.names
HeaderNames=['age','job','marital','education','default','housing','loan','contact','mo

for Idx in range(0,21):
    df_marketing_data=df_marketing_data.withColumnRenamed(DefColumnNames[Idx],HeaderNam
df_marketing_data = df_marketing_data.drop ('duration')
df_marketing_data.printSchema()
```

```
root
|-- age: integer (nullable = true)
|-- job: string (nullable = true)
|-- marital: string (nullable = true)
|-- education: string (nullable = true)
|-- default: string (nullable = true)
|-- housing: string (nullable = true)
|-- loan: string (nullable = true)
|-- contact: string (nullable = true)
|-- month: string (nullable = true)
|-- day_of_week: string (nullable = true)
|-- campaign: integer (nullable = true)
|-- pdays: integer (nullable = true)
|-- previous: integer (nullable = true)
|-- poutcome: string (nullable = true)
|-- emp_var_rate: double (nullable = true)
|-- cons_price_idx: double (nullable = true)
|-- cons_conf_idx: double (nullable = true)
|-- euribor3m: double (nullable = true)
|-- nr_employed: double (nullable = true)
|-- deposit: string (nullable = true)
```

In [12]:

```
# Segregação Variaveis Categorias e Numericas em Listas Específicas
categoricalColumns = []
numericCols = []
for i in df_marketing_data.dtypes:
    if i[1]=='string':
        categoricalColumns += [i[0]]
    elif i[1]=='int' or i[1]=='double':
        numericCols += [i[0]]

print(categoricalColumns)
print(numericCols)
```

```
['job', 'marital', 'education', 'default', 'housing', 'loan', 'contact',
'month', 'day_of_week', 'poutcome', 'deposit']
['age', 'campaign', 'pdays', 'previous', 'emp_var_rate', 'cons_price_id
x', 'cons_conf_idx', 'euribor3m', 'nr_employed']
```

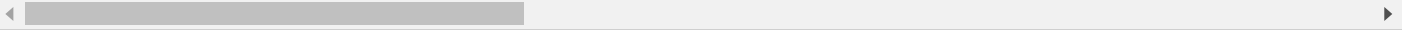
In [13]:

```
# Tratamento de Variáveis Categorias
stages = []
for categoricalCol in categoricalColumns:
    stringIndexer = StringIndexer(inputCol=categoricalCol, outputCol=categoricalCol+"Index")
    encoder = OneHotEncoder(inputCol=categoricalCol+"Index", outputCol=categoricalCol+"class")
    stages += [stringIndexer, encoder]

#numericCols = ['age', 'balance', 'duration', 'campaign', 'pdays', 'previous']
label_stringIdx = StringIndexer(inputCol = "deposit", outputCol = "label")
stages += [label_stringIdx]
```

In [14]:

```
## Assembler Inputs
assemblerInputs = ['jobclassVec', 'maritalclassVec', 'educationclassVec', 'defaultclassVec']
assembler = VectorAssembler(inputCols=assemblerInputs, outputCol="features")
stages += [assembler]
```



In [15]:

```
## Pipeline
pipeline = Pipeline(stages=stages)
pipelineModel = pipeline.fit(df_marketing_data)
df_marketing_data_prep = pipelineModel.transform(df_marketing_data)
df_marketing_data_prep.printSchema()
```

```
root
|-- age: integer (nullable = true)
|-- job: string (nullable = true)
|-- marital: string (nullable = true)
|-- education: string (nullable = true)
|-- default: string (nullable = true)
|-- housing: string (nullable = true)
|-- loan: string (nullable = true)
|-- contact: string (nullable = true)
|-- month: string (nullable = true)
|-- day_of_week: string (nullable = true)
|-- campaign: integer (nullable = true)
|-- pdays: integer (nullable = true)
|-- previous: integer (nullable = true)
|-- poutcome: string (nullable = true)
|-- emp_var_rate: double (nullable = true)
|-- cons_price_idx: double (nullable = true)
|-- cons_conf_idx: double (nullable = true)
|-- euribor3m: double (nullable = true)
|-- nr_employed: double (nullable = true)
|-- deposit: string (nullable = true)
|-- jobIndex: double (nullable = true)
|-- jobclassVec: vector (nullable = true)
|-- maritalIndex: double (nullable = true)
|-- maritalclassVec: vector (nullable = true)
|-- educationIndex: double (nullable = true)
|-- educationclassVec: vector (nullable = true)
|-- defaultIndex: double (nullable = true)
|-- defaultclassVec: vector (nullable = true)
|-- housingIndex: double (nullable = true)
|-- housingclassVec: vector (nullable = true)
|-- loanIndex: double (nullable = true)
|-- loanclassVec: vector (nullable = true)
|-- contactIndex: double (nullable = true)
|-- contactclassVec: vector (nullable = true)
|-- monthIndex: double (nullable = true)
|-- monthclassVec: vector (nullable = true)
|-- day_of_weekIndex: double (nullable = true)
|-- day_of_weekclassVec: vector (nullable = true)
|-- poutcomeIndex: double (nullable = true)
|-- poutcomeclassVec: vector (nullable = true)
|-- depositIndex: double (nullable = true)
|-- depositclassVec: vector (nullable = true)
|-- label: double (nullable = true)
|-- features: vector (nullable = true)
```

In [16]:

```
df_marketing_data_prep.take(5)
```

Out[16]:

```

[Row(age=56, job='housemaid', marital='married', education='basic.4y', de
fault='no', housing='no', loan='no', contact='telephone', month='may', da
y_of_week='mon', campaign=1, pdays=999, previous=0, poutcome='nonexisten
t', emp_var_rate=1.1, cons_price_idx=93.994, cons_conf_idx=-36.4, euribor
3m=4.857, nr_employed=5191.0, deposit='no', jobIndex=8.0, jobclassVec=Spa
rseVector(11, {8: 1.0}), maritalIndex=0.0, maritalclassVec=SparseVector
(3, {0: 1.0}), educationIndex=4.0, educationclassVec=SparseVector(7, {4:
1.0}), defaultIndex=0.0, defaultclassVec=SparseVector(2, {0: 1.0}), housi
ngIndex=1.0, housingclassVec=SparseVector(2, {1: 1.0}), loanIndex=0.0, lo
anclassVec=SparseVector(2, {0: 1.0}), contactIndex=1.0, contactclassVec=S
parseVector(1, {}), monthIndex=0.0, monthclassVec=SparseVector(9, {0: 1.
0}), day_of_weekIndex=1.0, day_of_weekclassVec=SparseVector(4, {1: 1.0}),
poutcomeIndex=0.0, poutcomeclassVec=SparseVector(2, {0: 1.0}), depositInd
ex=0.0, depositclassVec=SparseVector(1, {0: 1.0}), label=0.0, features=Sp
arseVector(52, {8: 1.0, 11: 1.0, 18: 1.0, 21: 1.0, 24: 1.0, 25: 1.0, 28:
1.0, 38: 1.0, 41: 1.0, 43: 56.0, 44: 1.0, 45: 999.0, 47: 1.1, 48: 93.994,
49: -36.4, 50: 4.857, 51: 5191.0})),
Row(age=57, job='services', marital='married', education='high.school',
default='unknown', housing='no', loan='no', contact='telephone', month='m
ay', day_of_week='mon', campaign=1, pdays=999, previous=0, poutcome='none
existent', emp_var_rate=1.1, cons_price_idx=93.994, cons_conf_idx=-36.4, e
uribor3m=4.857, nr_employed=5191.0, deposit='no', jobIndex=3.0, jobclassV
ec=SparseVector(11, {3: 1.0}), maritalIndex=0.0, maritalclassVec=SparseV
ector(3, {0: 1.0}), educationIndex=1.0, educationclassVec=SparseVector(7,
{1: 1.0}), defaultIndex=1.0, defaultclassVec=SparseVector(2, {1: 1.0}), h
ousingIndex=1.0, housingclassVec=SparseVector(2, {1: 1.0}), loanIndex=0.
0, loanclassVec=SparseVector(2, {0: 1.0}), contactIndex=1.0, contactclass
Vec=SparseVector(1, {}), monthIndex=0.0, monthclassVec=SparseVector(9,
{0: 1.0}), day_of_weekIndex=1.0, day_of_weekclassVec=SparseVector(4, {1:
1.0}), poutcomeIndex=0.0, poutcomeclassVec=SparseVector(2, {0: 1.0}), dep
ositIndex=0.0, depositclassVec=SparseVector(1, {0: 1.0}), label=0.0, feat
ures=SparseVector(52, {3: 1.0, 11: 1.0, 15: 1.0, 22: 1.0, 24: 1.0, 25: 1.
0, 28: 1.0, 38: 1.0, 41: 1.0, 43: 57.0, 44: 1.0, 45: 999.0, 47: 1.1, 48:
93.994, 49: -36.4, 50: 4.857, 51: 5191.0})),
Row(age=37, job='services', marital='married', education='high.school',
default='no', housing='yes', loan='no', contact='telephone', month='may',
day_of_week='mon', campaign=1, pdays=999, previous=0, poutcome='nonexiste
nt', emp_var_rate=1.1, cons_price_idx=93.994, cons_conf_idx=-36.4, euribo
r3m=4.857, nr_employed=5191.0, deposit='no', jobIndex=3.0, jobclassVec=Sp
arseVector(11, {3: 1.0}), maritalIndex=0.0, maritalclassVec=SparseVector
(3, {0: 1.0}), educationIndex=1.0, educationclassVec=SparseVector(7, {1:
1.0}), defaultIndex=0.0, defaultclassVec=SparseVector(2, {0: 1.0}), housi
ngIndex=0.0, housingclassVec=SparseVector(2, {0: 1.0}), loanIndex=0.0, lo
anclassVec=SparseVector(2, {0: 1.0}), contactIndex=1.0, contactclassVec=S
parseVector(1, {}), monthIndex=0.0, monthclassVec=SparseVector(9, {0: 1.
0}), day_of_weekIndex=1.0, day_of_weekclassVec=SparseVector(4, {1: 1.0}),
poutcomeIndex=0.0, poutcomeclassVec=SparseVector(2, {0: 1.0}), depositInd
ex=0.0, depositclassVec=SparseVector(1, {0: 1.0}), label=0.0, features=Sp
arseVector(52, {3: 1.0, 11: 1.0, 15: 1.0, 21: 1.0, 23: 1.0, 25: 1.0, 28:
1.0, 38: 1.0, 41: 1.0, 43: 37.0, 44: 1.0, 45: 999.0, 47: 1.1, 48: 93.994,
49: -36.4, 50: 4.857, 51: 5191.0})),
Row(age=40, job='admin.', marital='married', education='basic.6y', defau
lt='no', housing='no', loan='no', contact='telephone', month='may', day_o
f_week='mon', campaign=1, pdays=999, previous=0, poutcome='nonexistent',

```

```
emp_var_rate=1.1, cons_price_idx=93.994, cons_conf_idx=-36.4, euribor3m=
4.857, nr_employed=5191.0, deposit='no', jobIndex=0.0, jobclassVec=Sparse
Vector(11, {0: 1.0}), maritalIndex=0.0, maritalclassVec=SparseVector(3,
{0: 1.0}), educationIndex=5.0, educationclassVec=SparseVector(7, {5: 1.
0}), defaultIndex=0.0, defaultclassVec=SparseVector(2, {0: 1.0}), housing
Index=1.0, housingclassVec=SparseVector(2, {1: 1.0}), loanIndex=0.0, loan
classVec=SparseVector(2, {0: 1.0}), contactIndex=1.0, contactclassVec=Spa
rseVector(1, {}), monthIndex=0.0, monthclassVec=SparseVector(9, {0: 1.
0}), day_of_weekIndex=1.0, day_of_weekclassVec=SparseVector(4, {1: 1.0}),
poutcomeIndex=0.0, poutcomeclassVec=SparseVector(2, {0: 1.0}), depositInd
ex=0.0, depositclassVec=SparseVector(1, {0: 1.0}), label=0.0, features=Sp
arseVector(52, {0: 1.0, 11: 1.0, 19: 1.0, 21: 1.0, 24: 1.0, 25: 1.0, 28:
1.0, 38: 1.0, 41: 1.0, 43: 40.0, 44: 1.0, 45: 999.0, 47: 1.1, 48: 93.994,
49: -36.4, 50: 4.857, 51: 5191.0})),
Row(age=56, job='services', marital='married', education='high.school',
default='no', housing='no', loan='yes', contact='telephone', month='may',
day_of_week='mon', campaign=1, pdays=999, previous=0, poutcome='nonexiste
nt', emp_var_rate=1.1, cons_price_idx=93.994, cons_conf_idx=-36.4, euribo
r3m=4.857, nr_employed=5191.0, deposit='no', jobIndex=3.0, jobclassVec=Sp
arseVector(11, {3: 1.0}), maritalIndex=0.0, maritalclassVec=SparseVector
(3, {0: 1.0}), educationIndex=1.0, educationclassVec=SparseVector(7, {1:
1.0}), defaultIndex=0.0, defaultclassVec=SparseVector(2, {0: 1.0}), housi
ngIndex=1.0, housingclassVec=SparseVector(2, {1: 1.0}), loanIndex=1.0, lo
anclassVec=SparseVector(2, {1: 1.0}), contactIndex=1.0, contactclassVec=S
parseVector(1, {}), monthIndex=0.0, monthclassVec=SparseVector(9, {0: 1.
0}), day_of_weekIndex=1.0, day_of_weekclassVec=SparseVector(4, {1: 1.0}),
poutcomeIndex=0.0, poutcomeclassVec=SparseVector(2, {0: 1.0}), depositInd
ex=0.0, depositclassVec=SparseVector(1, {0: 1.0}), label=0.0, features=Sp
arseVector(52, {3: 1.0, 11: 1.0, 15: 1.0, 21: 1.0, 24: 1.0, 26: 1.0, 28:
1.0, 38: 1.0, 41: 1.0, 43: 56.0, 44: 1.0, 45: 999.0, 47: 1.1, 48: 93.994,
49: -36.4, 50: 4.857, 51: 5191.0}))]
```

In [17]:

```
#Configurando o modelo para 100 iterações
modelo = LogisticRegression(labelCol='label',featuresCol="features",maxIter=100)
```

In [18]:

```
# Divisão Teste/Treino
(marketing_model_treino, marketing_model_teste) = df_marketing_data_prep.randomSplit([0
```

In [19]:

```
# Ajuste Modelo
modelo_treino = modelo.fit(marketing_model_treino)
```

In [20]:

```
print("Coefficients: " + str(modelo_treino.coefficients))  
print("Intercept: " + str(modelo_treino.intercept))
```

```
Coefficients: [0.07228744920214636, -0.1823426344642905, 0.0430586707794122  
16, -0.01883575734932906, 0.020471150492064376, 0.309900243321121, 0.15349816  
856792098, -0.06432047069863493, -0.1400150462940802, -0.002065678868611054  
7, 0.24121218200088926, -0.016954238633392234, 0.02325671326311776, -0.042225  
64404035784, 0.054928930286601604, -0.08017466339569781, -0.0359619276043454  
05, 0.005322309614727278, -0.01938181814730077, 0.07707091629458984, 0.073784  
4541872536, 0.12767177016763268, -0.12429696421317507, -0.00685663740998846  
8, 0.02712818241933672, 0.014291818446949078, 0.022766872495474766, 0.4456242  
753007335, -0.5384796111547651, 0.34323584560280646, 0.07928836381942911, 0.2  
193636673715537, -0.1788984376052925, 0.10988741726505098, 0.109305924514261  
66, -0.09864335736514403, 1.2179087053600615, 0.039407318052363, -0.234098438  
02890684, 0.13166500396570288, 0.022384823939739673, -0.037130229052209335, -  
0.5670373053301918, 0.0005058246587037442, -0.043753943610546873, -0.0014133  
142529665966, -0.018416573586666075, -0.17781047722186252, 0.319901639609342  
6, 0.024852701484729234, -0.14656094785240206, -0.00533649993525627]  
Intercept: -2.0689878211161257
```

In [21]:

```
## Sumario  
print(modelo_treino.summary.predictions.stat)
```

<pyspark.sql.dataframe.DataFrameStatFunctions object at 0x7f12dc7a0f28>

In [22]:

```
# Salva o Modelo no HDFS  
hdfs_path = "/user/labdata/modelo_LR"  
modelo_treino.write().overwrite().save(hdfs_path)
```

In [23]:

```
# Regressão Logística  
modelo_treino2 = LogisticRegressionModel.load(hdfs_path)
```

In [24]:

```
## Executa a Predição  
predict = modelo_treino2.transform(marketing_model_treino)
```


In [25]:

```
## Exibe Predição
predict.show()
```

```
+---+-----+-----+-----+-----+-----+-----+-----+-----+
--+-----+-----+-----+-----+-----+-----+-----+-----+
-----+-----+-----+-----+-----+-----+-----+-----+
-----+-----+-----+-----+-----+-----+-----+-----+
-----+-----+-----+-----+-----+-----+-----+-----+
-----+-----+-----+-----+-----+-----+-----+-----+
-----+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+-----+-----+
-----+-----+
|age|      job|marital|  education|default|housing|loan|  contact|mon
th|day_of_week|campaign|pdays|previous|  poutcome|emp_var_rate|cons_pr
ice_idx|cons_conf_idx|euribor3m|nr_employed|deposit|jobIndex|  jobcla
ssVec|maritalIndex|maritalclassVec|educationIndex|educationclassVec|def
aultIndex|defaultclassVec|housingIndex|housingclassVec|loanIndex|  loanc
lassVec|contactIndex|contactclassVec|monthIndex|monthclassVec|day_of_we
ekIndex|day_of_weekclassVec|poutcomeIndex|poutcomeclassVec|depositIndex
|depositclassVec|label|      features|      rawPrediction|
probability|prediction|
+---+-----+-----+-----+-----+-----+-----+-----+-----+
```

In [26]:

```
results = predict.select(['probability', 'label'])
```

In [27]:

```
results_collect = results.collect()
results_list = [(float(i[0][0]), 1.0-float(i[1])) for i in results_collect]
scoreAndLabels = sc.parallelize(results_list)
```

In [28]:

```
# Exibe ROC
metrics = metric(scoreAndLabels)
print("The ROC score is (@itarações =100): ", metrics.areaUnderROC)
```

The ROC score is (@itarações =100): 0.7935651781484536

In [29]:


```
from sklearn.metrics import roc_curve, auc
from matplotlib import pyplot as plt

fpr = dict()
tpr = dict()
roc_auc = dict()

y_test = [i[1] for i in results_list]
y_score = [i[0] for i in results_list]

fpr, tpr, _ = roc_curve(y_test, y_score)
roc_auc = auc(fpr, tpr)

get_ipython().run_line_magic('matplotlib', 'inline')
plt.figure()
plt.plot(fpr, tpr, label='ROC curve (area = %0.2f)' % roc_auc)
plt.plot([0, 1], [0, 1], 'k--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver operating characteristic example')
plt.legend(loc="lower right")
plt.show()
display()
```



In []: